

Yi Hahn Pang

+65 8835 7353 | yihahn.pang@gmail.com | linkedin.com/in/yihahn-pang | github.com/yihahn-pang | yihahnpang.com

EDUCATION

Georgia Institute of Technology

Master of Science in Computer Science (OMSCS)

Atlanta, GA

Expected Aug. 2026 – Present

National University of Singapore

Bachelor of Engineering (Hons), Environmental Engineering & Data Analytics

Singapore

Aug. 2022 – June 2026

- **Academics:** Cumulative GPA: 4.46/5.0 (Distinction) | ASEAN Scholarship Recipient
- **Leadership & Athletics:** Won 7 athletic medals representing NUS College; Co-organized a 2-week symposium hosting 150 students from 10 ASEAN countries; Launched campus-wide BYOC waste-reduction campaign.

EXPERIENCE

Data Engineer Intern

Climate Impact X (CIX)

Jan. 2026 – May 2026

Singapore

- **Invited to return for a second term** based on technical ownership and effective communication.
- Architected the migration of legacy SQL pipelines to dbt, establishing a materialization strategy that reduced Snowflake compute time by 86% (from 66 to 9 mins).
- Enhanced pipeline reliability by developing comprehensive dbt tests, integrating custom Microsoft Teams alerts to surface business logic warnings without disrupting Airflow runs.
- Spearheaded AI integration by deploying a Microsoft Fabric Data Agent (FDA) within M365, utilizing Snowflake change tracking for automated data refreshes and natural language querying.

Data Engineer Intern

Climate Impact X (CIX)

Jan. 2025 – June 2025

Singapore

- Built 18 automated ETL pipelines (Airflow, PySpark, Azure) to ingest data from multiple sources into Snowflake and MongoDB, saving 10+ manual hours weekly.
- Refactored Snowflake schema by consolidating 40+ views into a Medallion Architecture, significantly improving data lineage and eliminating schema redundancies.
- Engineered 2 CI/CD pipelines (Bash, Azure DevOps) to automate Azure deployments and ensure strict environment parity.

Environmental Engineer Intern

Carbon GPT

May 2024 – Aug. 2024

Singapore

- Developed 4 LLM prototypes (Streamlit, Groq API) to automate manual processes; prototypes were subsequently utilized in successful sales pitches to UOB, CIMB, and Bursa.
- Engineered an ETL pipeline (Python, PostgreSQL) to process and optimize over 1.5M+ records, ensuring robust database scalability for downstream analysis.
- Led client engagements to gather technical requirements, ensuring product alignment with direct user feedback.

PROJECTS

Local Snowflake SQL Transpiler | *DuckDB, sqlglot, Python*

- Developed a local Snowflake SQL execution environment allowing developers to conduct fast, cost-free query testing and ensuring cross-database compatibility.

Real-Time Stock Streaming Dashboard | *Kafka, Streamlit*

- Developed a real-time stock trading simulator utilizing event streaming to visualize high-frequency market data.

Weather Analytics ELT Pipeline | *AWS Lambda, S3, Redshift*

- Built a fully cloud-native ELT pipeline to extract, load, and analyze historical Singapore weather trends.

SGX Derivatives Data Downloader CLI | *Python*

- Built a Python CLI tool to automate the retrieval of SGX derivatives data, complete with automated retry mechanisms, logging, and job scheduling support.

TECHNICAL SKILLS

Languages: Python (Pandas, NumPy, FastAPI), SQL, Java, R

Data Stack: Apache Spark (PySpark), Apache Airflow, Kafka, dbt, Snowflake, Databricks, PostgreSQL, MongoDB

Cloud, ML & BI: AWS, Azure, Microsoft Fabric, Docker, Git, CI/CD, PyTorch, Scikit-learn, Power BI

Certifications: AWS Certified Cloud Practitioner, Databricks Certified Data Engineer Associate